

```
# this cell is just auto formatting so I can be lazy and still have pretty code
#!pip install black[jupyter]
```

```
# from google.colab import drive
# drive.mount("/content/drive")
```

```
#!black /content/drive/MyDrive/'Colab Notebooks'/'DATA602_HW2.ipynb'
```

```
reformatted /content/drive/MyDrive/Colab Notebooks/DATA602_HW2.ipynb
```

```
All done! ✨ 🍰 ✨
1 file reformatted.
```

```
# install the necessary modules
#!pip install requests
#!pip install pandas
#!pip install numpy
#!pip install beautifulsoup4
```

```
# import the necessary modules
import requests
import pandas as pd
import numpy as np
from bs4 import BeautifulSoup
```

```
# use requests to pull the data
url = (
    "https://web.archive.org/web/20201112015618/https://www.spaceweatherlive"
    ".com/en/solar-activity/top-50-solar-flares.html"
)
```

```
r = requests.get(url)
```

```
# use beautifulsoup to parse the data
soup = BeautifulSoup(r.content, "html")
```

```
# look for the table we need with prettify
print(soup.prettify())
```

Show hidden output

```
# grab the table data
tb = soup.find("table")
```

```
# read the spaceweather live table into a dataframe
spwl_df = pd.read_html(str(tb))[0]
```

```
# change the column names
```

```
spwl_df.columns = [
    "rank",
    "x_class",
    "date",
    "region",
    "start_time",
    "max_time",
    "end_time",
    "movie",
]
```

```
# look at the top 5 in the dataframe
print(spwl_df.head())
```

```
# this completes step 1 of the assignment
```

	rank	x_class	date	region	start_time	max_time	end_time	\
0	1	X28+	2003/11/04	486	19:29	19:53	20:06	
1	2	X20+	2001/04/02	9393	21:32	21:51	22:03	
2	3	X17.2+	2003/10/28	486	09:51	11:10	11:24	
3	4	X17+	2005/09/07	808	17:17	17:40	18:03	
4	5	X14.4	2001/04/15	9415	13:19	13:50	13:55	

	movie
0	MovieView archive
1	MovieView archive
2	MovieView archive

```

3  MovieView archive
4  MovieView archive

# step 2 of the assignment
# drop the movie column and confirm that it was dropped
del spwl_df["movie"]
# commented out since no longer needed
# print(spwl_df.head())

# import datetime
from datetime import datetime

# combine the 3 time variables with the date
# for each row in the dataset
for index, row in spwl_df.iterrows():
    # for each of the three time columns
    for time_col in ["start_time", "max_time", "end_time"]:
        # Combine date and time
        combined_datetime = f"{row['date']} {row[time_col]}"
        # Convert to datetime format
        datetime_obj = datetime.strptime(combined_datetime, "%Y/%m/%d %H:%M")
        # Update the DataFrame
        spwl_df.at[index, time_col] = datetime_obj

# convert the time time columns to datetime format
spwl_df["start_time"] = pd.to_datetime(spwl_df["start_time"])
spwl_df["max_time"] = pd.to_datetime(spwl_df["max_time"])
spwl_df["end_time"] = pd.to_datetime(spwl_df["end_time"])

# Rename the time columns to end with _datetime
spwl_df = spwl_df.rename(
    columns={
        "start_time": "start_datetime",
        "max_time": "max_datetime",
        "end_time": "end_datetime",
    }
)

# replace - in the region column with Nan
# commented out since no longer needed
# print(spwl_df[spwl_df["region"] == "-"]) # there wasn't any that existed in the data
spwl_df["region"] = spwl_df["region"].replace("-", np.nan)

print(spwl_df.head())

# this completes step 2 of the assignment

   rank x_class    date  region  start_datetime  max_datetime \
0     1   X28+  2003/11/04    486  2003-11-04 19:29:00  2003-11-04 19:53:00
1     2   X20+  2001/04/02   9393  2001-04-02 21:32:00  2001-04-02 21:51:00
2     3  X17.2+  2003/10/28    486  2003-10-28 09:51:00  2003-10-28 11:10:00
3     4   X17+  2005/09/07    808  2005-09-07 17:17:00  2005-09-07 17:40:00
4     5  X14.4  2001/04/15   9415  2001-04-15 13:19:00  2001-04-15 13:50:00

   end_datetime
0  2003-11-04 20:06:00
1  2001-04-02 22:03:00
2  2003-10-28 11:24:00
3  2005-09-07 18:03:00
4  2001-04-15 13:55:00

# start step 3
# use requests to pull the data
# http://www.hcbravo.org/IntroDataSci/misc/waves_type2.html is missing data
url = "https://cdaw.gsfc.nasa.gov/CME_list/radio/waves_type2.html"
r = requests.get(url)

# use BeautifulSoup to parse the data
soup = BeautifulSoup(r.content, "html")

# look for the table we need with prettify
# print(soup.prettify())

# grab the table data
tb = soup.find("pre")

```

```

# read the table into a dataframe
data_text = soup.find("pre").text.split("\n")[12:-2]

# for testing only - verifying that it pulled only the text we want
# print(data_text)

# Split each line of text into a list of values
data_rows = [line.split() for line in data_text if line]

# Extracting the required columns
data_processed = []
for row in data_rows:
    start_date = row[0]
    start_time = row[1]
    end_date = row[2]
    end_time = row[3]
    start_frequency = row[4]
    end_frequency = row[5]
    flare_location = row[6]
    flare_region = row[7]
    flare_classification = row[8]
    cme_date = row[9]
    cme_time = row[10]
    cme_angle = row[11]
    cme_width = row[12]
    cme_speed = row[13]
    plot = row[14]
    data_processed.append(
        [
            start_date,
            start_time,
            end_date,
            end_time,
            start_frequency,
            end_frequency,
            flare_location,
            flare_region,
            flare_classification,
            cme_date,
            cme_time,
            cme_angle,
            cme_width,
            cme_speed,
            plot,
        ]
    )

# Convert the processed data into a DataFrame
nasa_df = pd.DataFrame(
    data_processed,
    columns=[
        "start_date",
        "start_time",
        "end_date",
        "end_time",
        "start_frequency",
        "end_frequency",
        "flare_location",
        "flare_region",
        "flare_classification",
        "cme_date",
        "cme_time",
        "cme_angle",
        "cme_width",
        "cme_speed",
        "plot",
    ],
)

# Display the first few rows of the DataFrame
print(nasa_df.head(10))

# Display the number of rows and columns to verify all data is there
print(nasa_df.shape)

# this ends step 3 of the project

```

	start_date	start_time	end_date	end_time	start_frequency	end_frequency	\
0	1997/04/01	14:00	04/01	14:15	8000	4000	
1	1997/04/07	14:30	04/07	17:30	11000	1000	
2	1997/05/12	05:15	05/14	16:00	12000	80	
3	1997/05/21	20:20	05/21	22:00	5000	500	
4	1997/09/23	21:53	09/23	22:16	6000	2000	
5	1997/11/03	05:15	11/03	12:00	14000	250	
6	1997/11/03	10:30	11/03	11:30	14000	5000	
7	1997/11/04	06:00	11/05	04:30	14000	100	
8	1997/11/06	12:20	11/07	08:30	14000	100	
9	1997/11/27	13:30	11/27	14:00	14000	7000	

	flare_location	flare_region	flare_classification	cme_date	cme_time	\
0	S25E16	8026		M1.3	04/01	15:18
1	S28E19	8027		C6.8	04/07	14:27
2	N21W08	8038		C1.3	05/12	05:30
3	N05W12	8040		M1.3	05/21	21:00
4	S29E25	8088		C1.4	09/23	22:02
5	S20W13	8100		C8.6	11/03	05:28
6	S16W21	8100		M4.2	11/03	11:11
7	S14W33	8100		X2.1	11/04	06:10
8	S18W63	8100		X9.4	11/06	12:10
9	N17E63	8113		X2.6	11/27	13:56

	cme_angle	cme_width	cme_speed	plot
0	74	79	312	PHTX
1	Halo	360	878	PHTX
2	Halo	360	464	PHTX
3	263	165	296	PHTX
4	133	155	712	PHTX
5	240	109	227	PHTX
6	233	122	352	PHTX
7	Halo	360	785	PHTX
8	Halo	360	1556	PHTX
9	98	91	441	PHTX

(522, 15)

```
# this starts step 4 of the project
```

```
# for testing/research purposes
```

```
# what are all the column datatypes - objects
```

```
# print(nasa_df.dtypes)
```

```
# START_DATE
```

```
# lets convert start_date to a date
```

```
nasa_df["start_date"] = pd.to_datetime(nasa_df["start_date"], errors="coerce")
```

```
# for testing/research purposes
```

```
# are all start dates good to go? Returns true so it is
```

```
# print(nasa_df['start_date'].notna().all())
```

```
# START_TIME
```

```
# going to use regex to check that all times are displayed correct
```

```
st_time_pat = r"^[01]?[0-9]|2[0-3]):[0-5][0-9]$"
```

```
# for testing/research purposes
```

```
# returns True so all start_time values are good to go
```

```
# print(nasa_df['start_time'].str.match(st_time_pat).all())
```

```
# END_DATE
```

```
# using regex to check that all the dates are displayed correctly
```

```
end_dt_pat = r"^(0[1-9]|1[0-2])/((0[1-9]|[12][0-9]|3[01]))$"
```

```
# for testing/research purposes
```

```
# returns True so all end dates are valid
```

```
# print(nasa_df['end_date'].str.match(end_dt_pat).all())
```

```
# END_TIME
```

```
# we can use the same regex pattern from start time to check end_time
```

```
# for testing/research purposes
```

```
# this came up false so lets try to find the row/rows that are causing the issue
```

```
# print(nasa_df['end_time'].str.match(st_time_pat).all())
```

```
is_valid = nasa_df["end_time"].str.match(st_time_pat)
```

```

invalid_rows = nasa_df[~is_valid]
# for testing/research purposes
# print(invalid_rows)
# the invalid end_times are 24:00. I can convert these to 23:59 with minimal
# change in the data
nasa_df["end_time"] = nasa_df["end_time"].replace("24:00", "23:59")

# START_FREQUENCY
# are there any non numerical start_frequency?
non_numerical_rows = nasa_df[
    nasa_df["start_frequency"].apply(lambda x: pd.to_numeric(x, errors="coerce")).isna()
]
# for testing purposes only - seeing what the rows look like
# print(non_numerical_rows)

# so the only non numerical rows for start_frequency
# is when start_frequency = '????'
nasa_df["start_frequency"] = nasa_df["start_frequency"].replace("????", np.nan)

# for testing purposes only - to confirm the change
# print(nasa_df[nasa_df['start_frequency'].isna()])

# END_FREQUENCY
# are there any non numerical end_frequency?
non_numerical_rows = nasa_df[
    nasa_df["end_frequency"].apply(lambda x: pd.to_numeric(x, errors="coerce")).isna()
]
# for testing purposes only - seeing what the rows look like
# print(non_numerical_rows)

# so the only non numerical rows for
# end_frequency is when end_frequency = '????'
nasa_df["end_frequency"] = nasa_df["end_frequency"].replace("????", np.nan)

# for testing purposes only - to confirm the change
# print(nasa_df[nasa_df['end_frequency'].isna()])

# FLARE_LOCATION
# flare location Back? is NaN but lets check for anything not in a format of
# (1 letter from N,E,S,W) + (0 or more numbers) +
# (0 or more letters from N,E,S,W) + (0 or more numbers) + (0 or 1 'b')

# first we need to strip flare_location of any trailing or leading zeros
nasa_df["flare_location"] = nasa_df["flare_location"].str.strip()

# Regex pattern
pattern = r"^[NESW]\d*[NESW]*\d*b?$"

# Find rows where 'flare_location' does not match the regex pattern
non_matching_rows = nasa_df[~nasa_df["flare_location"].str.match(pattern)]
# for testing purposes only - see what the non matching rows look like
# print(non_matching_rows)
# for testing purposes only - get the count of non matching rows (32)
# non_matching_count = (~nasa_df['flare_location'].str.match(pattern)).sum()
# print(non_matching_count)

# convert the rows to NaN
nasa_df["flare_location"] = np.where(
    nasa_df["flare_location"].str.match(pattern), nasa_df["flare_location"], np.nan
)

# for testing purposes only -
# confirm the same number of rows were converted (32)
# print(nasa_df[nasa_df['flare_location'].isna()])
# nan_count = nasa_df['flare_location'].isna().sum()
# print(nan_count)

# FLARE_REGION
# NOAA active region numbers should just be 4 to 5 numbers

```

```

# They also have things like
#     FILA (filament) and DSF(disappearing solar filament)
# but unless we need it later, we're going to mark them as NaN
# a filament is a loop burst thing that occurs on the sun
# looking at the data there is also
#     'EP', 'EP?', 'altr', and various lengths of '-----'
# pd.set_option("display.max_rows", None)
pattern = r"^\d+$"
non_matching_rows = nasa_df[~nasa_df["flare_region"].str.match(pattern)]
# for testing purposes -
#     looking at what all is in the column that doesn't fit the pattern
# print(non_matching_rows['flare_region'])
# pd.reset_option("display.max_rows")
# for testing purposes only - get the count of non matching rows (100)
# non_matching_count = (~nasa_df['flare_region'].str.match(pattern)).sum()
# print(non_matching_count)

# convert the rows to NaN
nasa_df["flare_region"] = np.where(
    nasa_df["flare_region"].str.match(pattern), nasa_df["flare_region"], np.nan
)

# for testing purposes only -
#     confirm the same number of rows were converted (100)
# nan_count = nasa_df['flare_region'].isna().sum()
# print(nan_count)

# FLARE_CLASSIFICATION
# for testing purposes only - reviewing what all flare_classification can be
# looks like it is always (letter)+(number)+(.)+(number OR blank)
# for empty ones it is always (----)
# pd.set_option("display.max_rows", None)
# print(nasa_df.groupby('flare_classification').size())
# pd.reset_option("display.max_rows")

# replace the flare_classifications when they are ----
nasa_df["flare_classification"] = nasa_df["flare_classification"].replace(
    "----", np.nan
)

# checking the numbers line up. 104 are now NaN
# nan_count = nasa_df['flare_classification'].isna().sum()
# print(nan_count)

# CME_DATE
cme_pat = r"^\d{2}/\d{2}$"

# returns False so there are some records that are not in the right format for cme_date
# print(nasa_df["cme_date"].str.match(cme_pat).all())
non_matching_rows = nasa_df[~nasa_df["cme_date"].str.match(cme_pat)]
# looks like cme_date is --/-- when it is empty and there is 20 of them
# print(non_matching_rows.groupby("cme_date").size())

# replace the cme_dates when they are --/--
nasa_df["cme_date"] = nasa_df["cme_date"].replace("--/--", np.nan)
# checking the numbers line up. 20 are now NaN
nan_count = nasa_df["cme_date"].isna().sum()
# print(nan_count)

# CME_TIME
cme_t_pat = r"^(?:[01]?[0-9]|2[0-3]):[0-5][0-9]|24:00$"

# returns False so there are some records that are not in the right format for cme_time
# print(nasa_df["cme_time"].str.match(cme_t_pat).all())
non_matching_rows = nasa_df[~nasa_df["cme_time"].str.match(cme_t_pat)]
# looks like cme_time is --:-- when it is empty and there is 20 of them
# print(non_matching_rows.groupby("cme_time").size())

# replace the cme_times when they are --:--
nasa_df["cme_time"] = nasa_df["cme_time"].replace("--:--", np.nan)
# checking the numbers line up. 20 are now NaN so it's good to go
nan_count = nasa_df["cme_time"].isna().sum()
# print(nan_count)

```

```

# CME_ANGLE
# The CPA column (cme_angle) contains angles in degrees for most rows, except for halo
# flares, which are coded as Halo. Create a new column that indicates if a row
# corresponds to a halo flare or not, and then replace Halo entries in the
# cme_angle column as NA.
cme_ang_pat = r"^(?:[0-9]{1,2}|[12][0-9]{2}|3[0-5][0-9]|360|Halo)$"

# returns False so there are records that don't meet our format
# print(nasa_df["cme_angle"].str.match(cme_ang_pat).all())
non_matching_rows = nasa_df[~nasa_df["cme_angle"].str.match(cme_ang_pat)]
# looks like cme_angle is ---- when it is empty and there is 21 of them
# print(non_matching_rows.groupby("cme_angle").size())
non_matching_rows

# replace the cme_angle when it is ----
nasa_df["cme_angle"] = nasa_df["cme_angle"].replace("----", np.nan)
# checking the numbers line up. 21 are now NaN so it's good to go
nan_count = nasa_df["cme_angle"].isna().sum()
# print(nan_count)

# create the new column as a True/False value for when the angle is 'Halo'
nasa_df["is_halo"] = nasa_df["cme_angle"] == "Halo"

# replace 'Halo' as NaN in cme_angle. Should be 264 of them for 285 NaN total now
nasa_df["cme_angle"] = nasa_df["cme_angle"].replace("Halo", np.nan)

# CME_WIDTH
# The width column indicates if the given value is a lower bound. Create a new column
# that indicates if width is given as a lower bound, and remove any non-numeric
# part of the width column.
cme_wid_pat = r"^(?:[0-9]{1,2}|[12][0-9]{2}|3[0-5][0-9]|360)$"

# returns False so there are records that don't meet our format
# print(nasa_df["cme_width"].str.match(cme_wid_pat).all())
non_matching_rows = nasa_df[~nasa_df["cme_width"].str.match(cme_wid_pat)]
# looks like cme_angle is ---- when it is empty and there is 21 of them
# print(non_matching_rows.groupby("cme_width").size())
# there is 4 rows that are --- , 16 rows that are ----, and one row that is 360h
# we also have 31 rows that indicate they are a lower bound since they start with >

# lets start by converting the --- rows
nasa_df["cme_width"] = nasa_df["cme_width"].replace("---", np.nan)
# converting the ---- rows
nasa_df["cme_width"] = nasa_df["cme_width"].replace("----", np.nan)
# converting the 360h row
nasa_df["cme_width"] = nasa_df["cme_width"].replace("360h", "360")

# make a flag indicating if cme_width is a lower bound or not
nasa_df["width_lower_bound"] = nasa_df["cme_width"].str.startswith(">")

# replace the > in cme_width
nasa_df["cme_width"] = nasa_df["cme_width"].str.replace(">", "", regex=False)

# CME_SPEED
cme_speed_pat = r"^\d+$"

# returns False so there are records that don't meet our format
# print(nasa_df["cme_speed"].str.match(cme_speed_pat).all())
non_matching_rows = nasa_df[~nasa_df["cme_speed"].str.match(cme_speed_pat)]
# looks like cme_speed is ---- when it is empty and there is 20 of them
# print(non_matching_rows.groupby("cme_speed").size())

# converting the ---- rows
nasa_df["cme_speed"] = nasa_df["cme_speed"].replace("----", np.nan)
# checking the numbers line up. 20 are now NaN so it's good to go
nan_count = nasa_df["cme_speed"].isna().sum()
# print(nan_count)

# PLOT
# all of the PLOT column is PHTX so there isn't any data needing cleaning
# print(nasa_df.groupby("plot").size())

# make the start_datetime flag by combining start_date and start_time
nasa_df["start_datetime"] = pd.to_datetime(

```

```

    nasa_df["start_date"].astype(str) + " " + nasa_df["start_time"]
)

# make the end_datetime flag by combining the year from start_date, end_date
# and the end_time flags
# Extract the year from 'start_date'
nasa_df["year"] = pd.to_datetime(nasa_df["start_date"]).dt.year.astype(str)

# Combine year with 'end_date', then combine with 'end_time'
# and convert to datetime format
nasa_df["end_datetime"] = pd.to_datetime(
    nasa_df["year"] + "-" + nasa_df["end_date"] + " " + nasa_df["end_time"]
)

# Combine year with 'cme_date', then combine with 'cme_time' and
# convert to datetime format
nasa_df["cme_datetime"] = pd.to_datetime(
    nasa_df["year"] + "-" + nasa_df["cme_date"] + " " + nasa_df["cme_time"]
)

# Drop the intermediate 'year' column
nasa_df.drop(columns="year", inplace=True)

# in order to make my dataset look exactly like the example in step 4
# convert flare_classification name to importance
nasa_df.rename(columns={"flare_classification": "importance"}, inplace=True)
# convert cme_angle name to cpa
nasa_df.rename(columns={"cme_angle": "cpa"}, inplace=True)
# convert cme_width name to width
nasa_df.rename(columns={"cme_width": "width"}, inplace=True)
# convert cme_speed name to speed
nasa_df.rename(columns={"cme_speed": "speed"}, inplace=True)
# drop start_date, start_time, end_date, end_time, cme_date, cme_time
nasa_df.drop(
    columns=[
        "start_date",
        "start_time",
        "end_date",
        "end_time",
        "cme_date",
        "cme_time",
    ],
    inplace=True,
)

# reorder the column so it displays the same
cols = (
    ["start_datetime"]
    + ["end_datetime"]
    + ["start_frequency"]
    + ["end_frequency"]
    + ["flare_location"]
    + ["flare_region"]
    + ["importance"]
    + ["cme_datetime"]
    + ["cpa"]
    + ["width"]
    + ["speed"]
    + ["plot"]
    + ["is_halo"]
    + ["width_lower_bound"]
)
nasa_df = nasa_df[cols]

# show the tidied table
nasa_df
# this concludes step 4

```


	start_datetime	end_datetime	start_frequency	end_frequency	flare_location	flare_region	importance	cme_datetime	c
0	1997-04-01 14:00:00	1997-04-01 14:15:00	8000	4000	S25E16	8026	M1.3	1997-04-01 15:18:00	
1	1997-04-07 14:30:00	1997-04-07 17:30:00	11000	1000	S28E19	8027	C6.8	1997-04-07 14:27:00	N
2	1997-05-12 05:15:00	1997-05-14 16:00:00	12000	80	N21W08	8038	C1.3	1997-05-12 05:30:00	N
3	1997-05-21 20:20:00	1997-05-21 22:00:00	5000	500	N05W12	8040	M1.3	1997-05-21 21:00:00	2
4	1997-09-23 21:53:00	1997-09-23 22:16:00	6000	2000	S29E25	8088	C1.4	1997-09-23 22:02:00	1
...
517	2017-09-17 11:45:00	2017-09-17 12:35:00	16000	900	S08E170	NaN	NaN	2017-09-17 12:00:00	N
518	2017-10-18 05:48:00	2017-10-18 12:40:00	16000	400	S06E123	NaN	NaN	2017-10-18 08:00:00	
519	2019-05-03 23:52:00	2019-05-04 00:16:00	13000	2300	N12E82	12740	C1.0	2019-05-03 23:24:00	

```
# part 2: analysis
# question 1: can I replicate the top 50 solar flare table from
# spaceweatherlive.com using the nasa data(nasa_df)
```

```
# first we need to understand the importance flag
# we're given that X28 is the highest
# for all classes besides X (so the A,B,C,and M classes ) the number after the
# letter only goes up to 9
# for the X class though, it goes up to 28 at which point the sensors cut out.
# reference: https://science.nasa.gov/science-research/heliophysics/space-weather/
# solar-flares/what-is-a-solar-flare/
```

```
# so to get the top 50 flares we can get the importance columns starting with X
# and the top numbers after that
```

```
# get just the flares that are X class
# had to fill the NaN records so I wouldn't get an error for trying
# to use the string function
x_flares = nasa_df[nasa_df["importance"].fillna("").str.startswith("X")]
```

```
# Sort the flares based on the number after 'X'
sorted_x_flares = x_flares.sort_values(
    by="importance", key=lambda x: x.str[1:].astype(float), ascending=False
)
```

```
# Extract the top 50
top_50_flares = sorted_x_flares[:50]
```

```
print(top_50_flares[["importance", "start_datetime", "flare_region"]])
```

```
# we're missing spaceweatherlives number 4 that occurred on 2005/09/07
# lets look at that start_datetime specifically
filtered_rows = nasa_df[
    nasa_df["start_datetime"].dt.date == pd.to_datetime("2005/09/07").date()
]
print(filtered_rows)
```

```
# INTERESTING! Looks like spaceweatherlive is incorrect on their 4th category.
# It is not X17.7 but X1.7. They also keep cutting the region data down to 4 numbers
# when there is 5 number regions.
# Overall, I'd trust the NASA data over what spaceweatherlive.com has
# which means I trust our list more than theirs.
```

240	X28.	2003-11-04 20:00:00	10486
117	X20.	2001-04-02 22:05:00	9393
233	X17.	2003-10-28 11:10:00	10486
126	X14.	2001-04-15 14:05:00	9415
234	X10.	2003-10-29 20:55:00	10486
8	X9.4	1997-11-06 12:20:00	8100
514	X9.3	2017-09-06 12:05:00	12673
328	X9.0	2006-12-05 10:50:00	10930
237	X8.3	2003-11-02 17:30:00	10486
515	X8.3	2017-09-10 16:02:00	NaN
288	X7.1	2005-01-20 07:15:00	10720
359	X6.9	2011-08-09 08:20:00	11263
331	X6.5	2006-12-06 19:00:00	10930
317	X6.2	2005-09-09 19:45:00	10808
82	X5.7	2000-07-14 10:30:00	9077
121	X5.6	2001-04-06 19:35:00	9415

```

443      X4.9 2014-02-25 00:56:00 11990
193      X4.8 2002-07-23 00:50:00 10039
104      X4.0 2000-11-26 17:00:00 9236
239      X3.9 2003-11-03 10:00:00 10488
286      X3.8 2005-01-17 10:00:00 10720
222      X3.6 2003-05-28 01:00:00 10365
332      X3.4 2006-12-13 02:45:00 10930
160      X3.4 2001-12-28 20:35:00 9756
192      X3.3 2002-07-20 21:30:00 10039
404      X3.2 2013-05-14 01:16:00 11748
201      X3.1 2002-08-24 01:45:00 10069
403      X2.8 2013-05-13 16:15:00 11748
487      X2.7 2015-05-05 22:24:00 12339
19       X2.7 1998-05-06 08:25:00 8210
238      X2.7 2003-11-03 01:15:00 10488
284      X2.6 2005-01-15 23:00:00 10720
142      X2.6 2001-09-24 10:45:00 9632
9         X2.6 1997-11-27 13:30:00 8113
276      X2.5 2004-11-10 02:25:00 10696
123      X2.3 2001-04-10 05:24:00 9415
99        X2.3 2000-11-24 15:25:00 9236
73        X2.3 2000-06-06 15:20:00 9026
345       X2.2 2011-02-15 02:10:00 11158
318       X2.1 2005-09-10 21:45:00 10808
361       X2.1 2011-09-06 22:30:00 11283
420       X2.1 2013-10-25 15:08:00 11882
7         X2.1 1997-11-04 06:00:00 8100
98        X2.0 2000-11-24 05:10:00 9236
125       X2.0 2001-04-12 10:20:00 9415
274       X2.0 2004-11-07 16:25:00 10696
285       X2.0 2005-01-17 09:25:00 10720
102       X1.9 2000-11-25 19:00:00 9236

      start_datetime end_datetime start_frequency end_frequency \
316 2005-09-07 18:05:00 2005-09-08 12000 200

      flare_location flare_region importance cme_datetime cpa width speed \
316 S11E77 10808 X1.7 NaT NaN NaN NaN

      plot is_halo width_lower_bound
316 PHTX False NaN

```

```

# part 2 question 2: write a function that finds the best matching row
# in the NASA data for each of the top 50 solar flares in SpaceWeatherLive data

```

```

# we will use xclass to match
# we can also match by date but from reviewing the time data in both datasets, there
# is mismatches in time even when it's the same observation
# we can use the region and if the nasa datas region is 5 in length then we'll use the
# last 4 to compare.
# Using these steps, we should be able to find all the errors in spaceweatherlives data

```

```

# is there any regions that are less than 1000 in the nasa data? if not we can safely
# add a 1 in front of all of the regions in temp_sp_df that start with 0
# I probably should also be converting region since its a number in the spaceweatherlive
# data and I want to compare the two
nasa_df["flare_region"] = pd.to_numeric(nasa_df["flare_region"], errors="coerce")
# nasa_df[nasa_df["flare_region"] < 1000]
# there's no regions under 1000 in the nasa data so if the spaceweather live data says
# 999 or under, it really should be 10999 since spaceweatherlive cut off the leading 1
# the next question in this is are there regions above 11000
# nasa_df[nasa_df["flare_region"] > 11000]
# there is
# so this means we can not just add a 1 in front of every region in the spaceweather
# data that is under 1000 and think we have fixed all of their messed up data
# However, adding a 1 in front of every region under 1000 will fix it quite a bit
# it just wont fix things like region 11263 in the nasa data being 1263 in the
# space weather data
# convert region to being numerical and fix the records under 1000
spwl_df["region"] = pd.to_numeric(spwl_df["region"], errors="coerce")
spwl_df.loc[spwl_df["region"] < 1000, "region"] += 10000

```

```

# getting the two tables into a format where it's easy to make comparisons
# take just the columns we will use for comparison from the spaceweatherlive data
temp_sp_df = spwl_df[["rank", "x_class", "date", "region"]]
# get rid of the + in the x_class column since the nasa data doesn't do that
# I had to make a copy and set it equal to avoid a warning message
# normally, I wouldn't concern myself with a warning message but I'm erring on the
# side of caution since this is HW

```

```

temp_sp_df = temp_sp_df.copy()
temp_sp_df.loc[:, "x_class"] = temp_sp_df["x_class"].str.rstrip("+")

# take just the columns we will use for comparison from the nasa data
# rename the columns so they line up with the spaceweatherlive data
temp_nasa_df = nasa_df[["importance", "start_datetime", "flare_region"]].rename(
    columns={
        "importance": "x_class",
        "flare_region": "region",
        "start_datetime": "date",
    }
)

# need to change start_datetime in the temp_nasa_df to just have the date and no time
temp_nasa_df["date"] = temp_nasa_df["date"].dt.strftime("%Y/%m/%d")

# convert the nasa datasets region flag to numeric and get rid of the decimal point
temp_nasa_df["region"] = pd.to_numeric(temp_nasa_df["region"], errors="coerce")
# Convert to integer to remove decimal points
temp_nasa_df["region"].fillna(-1, inplace=True)
temp_nasa_df["region"] = temp_nasa_df["region"].astype(int)
temp_nasa_df["region"] = temp_nasa_df["region"].replace(-1, np.nan)

# make the x_class in each table end with a decimal and a 0 so that they match
# better
# Define a function to add a decimal point and a 0 if not present
def add_decimal_and_zero(x):
    # Convert x to string if it's not already a string
    x = str(x)

    if "." not in x:
        return x + ".0"
    elif x.endswith("."):
        return x + "0"
    return x

# Apply the function to the x_class column
temp_sp_df["x_class"] = temp_sp_df["x_class"].apply(add_decimal_and_zero)
temp_nasa_df["x_class"] = temp_nasa_df["x_class"].apply(add_decimal_and_zero)

# make a new column on the nasa data that will hold the space weather rank if it
# exists in the space weather data
# initially set them all to null
nasa_df["sp_weather_rank"] = np.nan

def find_closest_match(row, df):
    # Check for exact matches in x_class
    x_class_match = df["x_class"] == row["x_class"]

    # Compute the absolute difference in days for the date column
    date_diff = (pd.to_datetime(df["date"]) - pd.to_datetime(row["date"])).abs().dt.days

    # Compare region values
    if len(str(row["region"])) == 4:
        region_diff = df["region"].astype(str).str[-4:] != str(row["region"])
    else:
        region_diff = df["region"] != row["region"]

    # Combine the differences to get a total "distance"
    # We give a high penalty (e.g., 1000 days) for non-matching x_class and region
    total_distance = (
        (~x_class_match * 100) + (region_diff * 100) + (date_diff != 0) * 1000
    )

    # Find the index of the row with the smallest distance
    closest_idx = total_distance.idxmin()

    # Return the closest row
    return df.loc[closest_idx]

# Finally, use your function to add a new column to the NASA dataset indicating its

```

```

# rank according to SpaceWeatherLive, if it appears in that dataset.
# Loop over each row in temp_sp_df
for i in range(len(temp_sp_df)):
    row = temp_sp_df.iloc[i] # Get the row at index i from temp_sp_df
    closest_match = find_closest_match(row, temp_nasa_df)

    # Check if sp_weather_rank is non-blank for the closest_match
    existing_rank = nasa_df.loc[closest_match.name, "sp_weather_rank"]
    new_rank = int(temp_sp_df.iloc[i]["rank"])

    # Since some records in the nasa data match to two records in the spaceweatherlive
    # data, update the rank if it exists instead of overwriting it
    if pd.notna(existing_rank): # If non-blank
        nasa_df.loc[
            closest_match.name, "sp_weather_rank"
        ] = f"{existing_rank}/{new_rank}"
    else:
        nasa_df.loc[closest_match.name, "sp_weather_rank"] = new_rank

# pull just the nasa data that has a sp_weather_rank
filtered_nasa_df = nasa_df[nasa_df["sp_weather_rank"].notna()]
# print(filtered_nasa_df)

# Create a temporary column for sorting
# I had to make a copy and set it equal to avoid a warning message
# normally, I wouldn't concern myself with a warning message but I'm erring on the
# side of caution since this is HW
filtered_nasa_df = filtered_nasa_df.copy()
filtered_nasa_df["sp_weather_rank"] = filtered_nasa_df["sp_weather_rank"].astype(str)
filtered_nasa_df["temp_sort"] = (
    filtered_nasa_df["sp_weather_rank"].str.split("/").str[0].astype(float)
)

# Sort the DataFrame based on the temporary column
sorted_df = filtered_nasa_df.sort_values(by="temp_sort", ascending=True)

# Drop the temporary column if it's no longer needed
sorted_df.drop(columns=["temp_sort"], inplace=True)

# print out the results
print(sorted_df[["start_datetime", "flare_region", "importance", "sp_weather_rank"]])

# I think these results are pretty good. It's correctly pulling the mistake for rank 4
# I don't love that rank 15, 16 and 31 all match to the same nasa record but when I
# looked at the data, there really doesn't seem to be any better match
# From looking at the records that I know should match, and comparing their
# start, max and end times, I think we're more likely to introduce more false positives
# if we were to include those flags.

```

	start_datetime	flare_region	importance	sp_weather_rank
240	2003-11-04 20:00:00	10486.0	X28.	1.0
117	2001-04-02 22:05:00	9393.0	X20.	2.0
233	2003-10-28 11:10:00	10486.0	X17.	3.0
316	2005-09-07 18:05:00	10808.0	X1.7	4.0/20
126	2001-04-15 14:05:00	9415.0	X14.	5.0
234	2003-10-29 20:55:00	10486.0	X10.	6.0
8	1997-11-06 12:20:00	8100.0	X9.4	7.0
514	2017-09-06 12:05:00	12673.0	X9.3	8.0
328	2006-12-05 10:50:00	10930.0	X9.0	9.0
237	2003-11-02 17:30:00	10486.0	X8.3	10.0
515	2017-09-10 16:02:00	NaN	X8.3	11.0
288	2005-01-20 07:15:00	10720.0	X7.1	12.0
359	2011-08-09 08:20:00	11263.0	X6.9	13.0
331	2006-12-06 19:00:00	10930.0	X6.5	14.0
317	2005-09-09 19:45:00	10808.0	X6.2	15.0/16/31
82	2000-07-14 10:30:00	9077.0	X5.7	17
121	2001-04-06 19:35:00	9415.0	X5.6	18
375	2012-03-07 01:00:00	11429.0	X5.4	19
231	2003-10-26 07:00:00	10486.0	X1.2	21
135	2001-08-25 16:50:00	9591.0	X5.3	22
443	2014-02-25 00:56:00	11990.0	X4.9	23/24
193	2002-07-23 00:50:00	10039.0	X4.8	25
104	2000-11-26 17:00:00	9236.0	X4.0	26
239	2003-11-03 10:00:00	10488.0	X3.9	27/28
286	2005-01-17 10:00:00	10720.0	X3.8	29
0	1997-04-01 14:00:00	8026.0	M1.3	30
222	2003-05-28 01:00:00	10365.0	X3.6	32/33
332	2006-12-13 02:45:00	10930.0	X3.4	34
160	2001-12-28 20:35:00	9756.0	X3.4	35

192	2002-07-20 21:30:00	10039.0	X3.3	36/37/38
404	2013-05-14 01:16:00	11748.0	X3.2	39
201	2002-08-24 01:45:00	10069.0	X3.1	40/41
187	2002-07-15 21:15:00	10030.0	M1.8	42
403	2013-05-13 16:15:00	11748.0	X2.8	43/45
157	2001-12-11 12:45:00	NaN	NaN	44
487	2015-05-05 22:24:00	12339.0	X2.7	46
238	2003-11-03 01:15:00	10488.0	X2.7	47
19	1998-05-06 08:25:00	8210.0	X2.7	48
284	2005-01-15 23:00:00	10720.0	X2.6	49
142	2001-09-24 10:45:00	9632.0	X2.6	50

```
# part 2 question 3
# I'm curious to see if solar flares are more likely in certain months of the year
# so I'll make a plot where the x-axis is the months and the y-axis is the number of
# solar flares
# maybe I'll get fancy with it and I can make a z-axis that is the years so we can
# look at how different months look throughout the years
# Overall, it looks like we had more flares (or more data for them) before 2006
# Nothing is sticking out to say that a specific month has more flares than others
# 2001 may have had the most flares of any year
```


```
# import the necessary modules
import seaborn as sns
import matplotlib.pyplot as plt

# Extract month and year from start_datetime
nasa_df["month"] = nasa_df["start_datetime"].dt.month
nasa_df["year"] = nasa_df["start_datetime"].dt.year

# Group by month and year and count the number of records
grouped = nasa_df.groupby(["year", "month"]).size().reset_index(name="count")

# Pivot the data
heatmap_data = grouped.pivot(index="year", columns="month", values="count")

# Plot the heatmap
plt.figure(figsize=(12, 7))
sns.heatmap(heatmap_data, cmap="YlGnBu", annot=True, fmt="g")
plt.title("Number of Solar Flares by Month and Year")
plt.xlabel("Month")
plt.ylabel("Year")
plt.show()
```



```

#!pip install plotly

# import the necessary module
import plotly.graph_objects as go

# get the data we need
nasa_df["month"] = nasa_df["start_datetime"].dt.month
nasa_df["year"] = nasa_df["start_datetime"].dt.year
grouped = nasa_df.groupby(["year", "month"]).size().reset_index(name="count")

# Map the count values to a colorscale
norm = plt.Normalize(grouped["count"].min(), grouped["count"].max())
colors = plt.cm.viridis(norm(grouped["count"]))

# Create lines to simulate bars with colors based on count and increased width
lines = []
for i, row in grouped.iterrows():
    color = f"rgb({int(colors[i][0]*255)}, {int(colors[i][1]*255)}, \
    {int(colors[i][2]*255)})"
    lines.append(
        go.Scatter3d(
            x=[row["month"], row["month"]],
            y=[row["year"], row["year"]],
            z=[0, row["count"]],
            mode="lines",
            line=dict(color=color, width=20),
        )
    ) # Increased width to 20

# Determine the range of years to display
min_year = min(grouped["year"].min(), 1997)
max_year = max(grouped["year"].max(), 2020)

# Create the 3D bar chart
fig = go.Figure(data=lines)

# Set labels, title, hide the legend, and specify tick values
fig.update_layout(
    scene=dict(
        xaxis_title="Month",
        yaxis_title="Year",
        zaxis_title="Number of Flares",
        xaxis=dict(
            tickvals=list(range(1, 13)), ticktext=[str(i) for i in range(1, 13)]
        ), # Every month from 1 to 12
        yaxis=dict(
            tickvals=list(range(min_year, max_year + 1)),
            ticktext=[str(i) for i in range(min_year, max_year + 1)],
        ), # Every year from min_year to max_year
    ),
    title="Number of Solar Flares by Month and Year",
    showlegend=False, # Hide the legend
)

fig.show()

```

Number of Solar Flares by Month and Year

